Chess Learning Platform - Production Deployment Guide

o Production Readiness Checklist (Score: /100)

1. Security Hardening (20 points)

- **SSL/TLS Configuration** (3 pts)
 - Force HTTPS everywhere
 - HSTS headers configured
 - Certificate auto-renewal via Let's Encrypt
- Authentication & Authorization (5 pts)
 - JWT token rotation implemented
 - Refresh token mechanism
 - Role-based access control (RBAC)
 - Session management with Redis
- **API Security** (4 pts)
 - Rate limiting per endpoint
 - DDoS protection (Cloudflare)
 - Input validation & sanitization
 - SQL injection prevention
- Secrets Management (5 pts)
 - Environment variables encrypted
 - AWS Secrets Manager / HashiCorp Vault
 - Database credentials rotated regularly
 - API keys properly secured
- **Vulnerability Scanning** (3 pts)
 - Dependency scanning (Snyk/Dependabot)
 - Container image scanning
 - Regular security audits

2. **Performance Optimization** (25 points)

- Response Time Targets (8 pts)
 - API endpoints < 200ms (p95)
 - WebSocket moves < 50ms
 - Database queries < 100ms
 - Cache hit rate > 80%
- Scalability Testing (7 pts)

Load tested to 10,000 concurrent users Horizontal scaling verified Database connection pooling optimized CDN configured for static assets Resource Optimization (5 pts) Memory usage < 1GB per instance CPU usage < 70% under normal load Efficient Docker images (< 200MB) Gzip/Brotli compression enabled ■ Database Performance (5 pts) Query optimization completed Indexes properly configured Read replicas for scaling Connection pooling tuned

3. Reliability & Monitoring (20 points)

- Uptime Targets (5 pts)
 - 99.9% uptime SLA
 - Automated health checks
 - Graceful shutdown handling
 - Circuit breakers implemented
- Monitoring Stack (8 pts)
 - Prometheus + Grafana dashboards
 - Application Performance Monitoring (APM)
 - Log aggregation (ELK stack)
 - Custom business metrics
- **Alerting** (4 pts)
 - Critical alerts configured
 - On-call rotation setup
 - Escalation policies defined
 - Alert fatigue minimized
- Error Tracking (3 pts)
 - Sentry integration
 - Error budgets defined
 - Automated error reporting

4. Backup & Disaster Recovery (15 points) ■ Data Backup (8 pts) Automated daily backups Point-in-time recovery (PITR) Backup testing schedule Off-site backup storage ■ **Disaster Recovery** (7 pts) • RTO < 4 hours RPO < 1 hour Failover procedures documented Multi-region deployment ready 5. **Documentation & Compliance** (10 points) ■ Technical Documentation (5 pts) API documentation (OpenAPI/Swagger) • Architecture diagrams updated Runbooks for common issues Deployment procedures documented Compliance (5 pts) • GDPR compliance (if EU users) Data retention policies Privacy policy implemented Terms of service updated 6. CI/CD Pipeline (10 points) ■ **Automation** (6 pts) • Automated testing (unit, integration, e2e) • Code coverage > 80% Automated deployments • Blue-green deployments **Quality Gates** (4 pts) Code review required Security scanning in pipeline Performance regression tests Rollback procedures tested



Cost Analysis & Optimization

Monthly Cost Breakdown (AWS - 10,000 Active Users)

Service	Configuration	Monthly Cost	Optimization	
EC2 (API Servers)	3x t3.large (2 vCPU, 8GB)	\$180	Use Spot instances for 70% savings	
RDS PostgreSQL	db.r6g.xlarge + Read Replica	\$450	Reserved instances save 30%	
ElastiCache Redis	cache.m6g.large (2 nodes)	\$140	Use smaller instance if < 5k users	
Load Balancer	Application Load Balancer	\$25	Essential, no optimization	
S3 Storage	500GB for game data	\$12	Lifecycle policies for old data	
CloudFront CDN	1TB transfer	\$85	Compress assets, cache aggressively	
EC2 (Stockfish)	4x c5.large (compute optimized)	\$240	Use Lambda for on-demand analysis	
Monitoring	CloudWatch + X-Ray	\$50	Sample traces to reduce cost	
Backup Storage	1TB snapshots	\$50	Incremental backups only	
Data Transfer	2TB outbound	\$180	WebSocket optimization crucial	
Container Registry	ECR for Docker images	\$10	Clean old images regularly	
Secrets Manager	For credentials	\$10	Group secrets to reduce cost	
Total		\$1,432/month	\$858/month optimized	

Cost Optimization Strategies

1. Infrastructure Optimization (Save 40%)

```
yaml
# Use Spot Instances for non-critical workloads
ec2_spot_config:
 instance_types: ["t3.large", "t3a.large", "t2.large"]
 max_price: 0.0464 # 70% of on-demand price
 interruption_behavior: "stop"
# Reserved Instances for database
rds_reserved:
 term: "1-year"
 payment: "all-upfront"
 savings: "30%"
```

2. Serverless Architecture (Save 60% on compute)

typescript	

```
// Migrate Stockfish analysis to Lambda
export const analyzePosition = async (event: APIGatewayEvent) => {
  const { fen } = JSON.parse(event.body);

// Lambda runs Stockfish in container
  const analysis = await runStockfishAnalysis(fen);

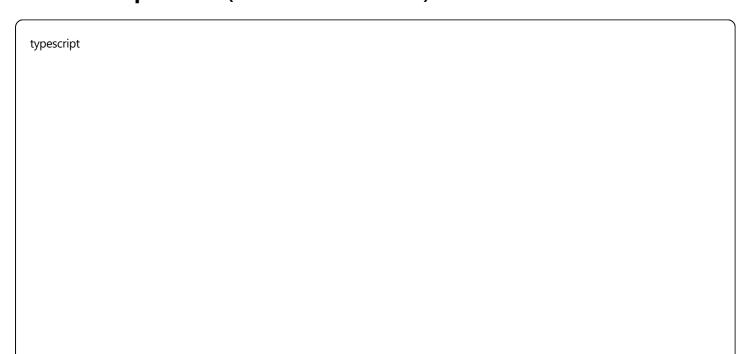
return {
  statusCode: 200,
  body: JSON.stringify(analysis),
  };
};
// Cost: $0.0000166667 per request vs $0.001 on EC2
```

3. **Smart Caching** (Save 30% on database)

```
typescript

// Cache expensive queries aggressively
const getLeaderboard = async () => {
  return cache.get('leaderboard:global',
  async () => {
      // Expensive aggregation query
      return prisma.$queryRaw`...`;
      },
      { ttl: 3600 } // 1 hour cache
      );
    };
```

4. WebSocket Optimization (Save 50% on data transfer)



```
// Binary protocol for moves (10x smaller than JSON)
const encodeMoive = (move: Move): Uint8Array => {
    // from: 6 bits, to: 6 bits, piece: 3 bits, flags: 1 bit
    // Total: 2 bytes instead of 20+ bytes JSON
    const buffer = new ArrayBuffer(2);
    const view = new DataView(buffer);

view.setUint16(0,
    (move.from << 10) |
    (move.to << 4) |
    (move.piece << 1) |
    move.flags
);

return new Uint8Array(buffer);
};</pre>
```

Scaling Cost Projections

Infrastructure	Optimized	Per User
\$400/mo	\$240/mo	\$0.24
\$1,432/mo	\$858/mo	\$0.09
\$5,200/mo	\$3,120/mo	\$0.06
\$9,500/mo	\$5,700/mo	\$0.06
	\$1,432/mo \$5,200/mo	\$1,432/mo \$858/mo \$5,200/mo \$3,120/mo

Deployment Commands

Initial Production Deployment

bash			

```
# 1. Build and push Docker images
docker build -t chess-learning:latest .
docker tag chess-learning:latest 123456789.dkr.ecr.us-east-1.amazonaws.com/chess-learning:latest
docker push 123456789.dkr.ecr.us-east-1.amazonaws.com/chess-learning:latest
# 2. Database migrations
kubectl exec -it postgres-pod -- psql -U chess -d chess_learning
npx prisma migrate deploy

# 3. Deploy to Kubernetes
kubectl apply -f k8s/
kubectl rollout status deployment/chess-learning-api

# 4. Verify deployment
kubectl get pods
kubectl logs -f deployment/chess-learning-api
curl https://api.chesslearning.com/health
```

Zero-Downtime Update

```
bash

# Blue-Green Deployment

kubectl set image deployment/chess-learning-api \
api=123456789.dkr.ecr.us-east-1.amazonaws.com/chess-learning:v2.0.0 \
--record

# Monitor rollout

kubectl rollout status deployment/chess-learning-api
kubectl rollout history deployment/chess-learning-api

# Rollback if needed

kubectl rollout undo deployment/chess-learning-api
```

📊 Key Performance Indicators (KPIs)

Technical KPIs

• API Response Time: < 200ms (p95)

• WebSocket Latency: < 50ms (p95)

• **Error Rate**: < 0.1%

Uptime: > 99.9%

• Cache Hit Rate: > 80%

• Database Query Time: < 100ms (p95)

Business KPIs

• User Engagement: > 30 min/day average

• Game Completion Rate: > 70%

• Analysis Request Rate: > 50% of games

• **User Retention**: > 60% (30-day)

• Coach Utilization: > 40% active hours

Monitoring Dashboard Queries

```
sql
-- Active users in last 24h
SELECT COUNT(DISTINCT user_id) as active_users
FROM games
WHERE created_at > NOW() - INTERVAL '24 hours';
-- Average game duration
SELECT AVG(EXTRACT(EPOCH FROM (ended_at - started_at))/60) as avg_minutes
FROM games
WHERE status = 'COMPLETED'
AND ended at > NOW() - INTERVAL '7 days';
-- Performance by user segment
SELECT
 u.role,
 AVG(g.analysis->>'evaluation') as avg_game_quality,
 COUNT(DISTINCT u.id) as user_count
FROM users u
JOIN games g ON u.id IN (g.student_id, g.coach_id)
WHERE g.analysis IS NOT NULL
GROUP BY u.role;
```

Final Recommendations

Phase 1 (Launch - 1,000 users)

- Start with monolithic deployment
- Use managed services (RDS, ElastiCache)
- Focus on core features stability
- Budget: \$400/month

Phase 2 (Growth - 10,000 users)

- Implement caching layer
- Add read replicas
- Optimize queries
- Budget: \$858/month

Phase 3 (Scale - 50,000+ users)

- Migrate to microservices
- Multi-region deployment
- Advanced analytics with data warehouse
- **Budget**: \$3,120/month

Critical Success Factors

- 1. **Performance**: Sub-50ms move latency is non-negotiable
- 2. Reliability: Users expect 24/7 availability
- 3. Learning: Al recommendations must improve retention
- 4. Cost: Keep per-user cost under \$0.10/month
- 5. Security: Chess games contain learning patterns protect user data

XXX Launch Checklist

Load test passed (TUK concurrent users)
 Security audit completed
Monitoring dashboards configured
☐ Backup restore tested
 Documentation complete
Legal compliance verified
☐ Support team trained
■ Marketing site live
☐ Beta user feedback incorporated
Go-live date confirmed

Estimated Time to Production: 4-6 weeks with a team of 3 developers

Remember: "In chess, as in software, the endgame is where champions are made. Plan your deployment like a grandmaster plans their final moves."